

# Car\_price\_Prediction

May 2, 2025

## 0.1 Problem Statement

The goal of this project is to build a robust machine learning model to predict the selling prices of used cars based on key vehicle attributes and historical sales data. The dataset, sourced from Mobil123 through web scraping, contains both numerical and categorical features that describe a car's specifications, usage, and ownership history.

The primary business problem is to enable more accurate, data-driven car valuations that benefit both buyers and sellers in the used car market. To address this, the project focuses on:

1. Identifying which features significantly influence car prices through exploratory data analysis and statistical testing.
2. Managing multicollinearity among highly correlated variables.
3. Selecting and engineering the most relevant features to improve model performance.
4. Comparing and evaluating multiple regression models (XGBoost, Gradient Boosting, Random Forest) to determine the most accurate and efficient approach for price prediction.

The project concludes with the deployment of the best-performing model (XGBoost), integrated with full preprocessing via a pipeline. This model is packaged and ready for real-world application through an interactive Streamlit-based web app.

## 0.2 Summary of Insights for Car Selling Price Prediction

Statistical analysis was conducted to understand the factors influencing car selling prices and prepare the data for predictive modeling.

---

### Key Findings from EDA and Correlation Analysis:

- **Strong Positive Drivers of Price:** Numerical features like `max_power`, `torque`, `engine size`, and the manufacturing `year` showed strong positive correlations with the selling price. This indicates that cars with higher specifications and newer models tend to command higher prices.
- **Negative Impact of Usage and Age:** Numerical features such as `car_age`, `mileage`, and `km_driven` exhibited negative correlations, suggesting that older and more used vehicles generally have lower selling prices.

**Categorical Feature Importance:** 1. Welch's ANOVA analysis indicated that categorical features like `fuel_type` and `transmission` have a statistically significant impact on the mean selling

price.

2. The number of previous owners (**owner**) also showed a moderate influence.
3. Despite the relatively low F-statistic, (**name**) is included in model training because it is a high-cardinality categorical feature (many unique car names) and is highly imbalanced. These characteristics dilute the overall variance captured by ANOVA yet the specific car make and model are strong determinants of car price.
4. **seller\_type** is considered less likely to contribute meaningfully to the model and was excluded. This is because it reflects the context of the transaction rather than intrinsic characteristics of the car.

---

### Multicollinearity Management:

1. A strong positive correlation has been identified between **engine**, **max\_power**, and **torque**.
2. A perfect negative correlation between **car\_age** and **year**.

One representative feature from each correlated group (**engine** or **max\_power**, and **year**) was chosen for model training to avoid redundancy.

3. Given that **mileage** and **km\_driven** represent the same concept of car usage, **km\_driven** has been used for the model training due to its slightly stronger negative correlation with price.
- **Feature Exclusion** **seller\_type** is considered less likely to contribute meaningfully to the model and was excluded. This is because it reflects the context of the transaction rather than intrinsic characteristics of the car.

---

### Model Evaluation:

Model	MAE	MSE	RMSE	R <sup>2</sup> Score
XGBoost (no outliers)	65,224.45	8,681,822,530.02	93,176.30	0.86
XGBoost	87,513.79	32,608,928,997.74	180,579.43	0.85
Gradient Boosting	88,424.50	36,732,389,083.97	191,656.96	0.83
Random Forest	91,316.73	39,347,158,917.14	198,361.18	0.82

XGBoost delivered the best performance across all evaluation metrics, making it the ideal model for predicting car prices. Removing outliers greatly improved its performance by cutting RMSE nearly in half and reducing MSE from 32.6 billion to 8.7 billion. This optimized version of XGBoost will be deployed as the final prediction tool.

```
[ ]: # Importing Libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from scipy.stats import levene
import statsmodels.api as sm
```

```

from statsmodels.formula.api import ols
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.impute import SimpleImputer
from sklearn.pipeline import Pipeline
from sklearn.compose import ColumnTransformer
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
from sklearn.ensemble import RandomForestRegressor
from sklearn.ensemble import GradientBoostingRegressor
from sklearn.model_selection import GridSearchCV
import xgboost as xgb
from xgboost import XGBRegressor
import pickle
import joblib
from matplotlib.ticker import ScalarFormatter

import warnings
warnings.filterwarnings('ignore')

```

```
[ ]: pd.set_option('display.float_format', '{:,.2f}'.format)
```

```

[ ]: # Reading the data from Google sheets
from google.colab import auth
auth.authenticate_user()
import gspread
from google.auth import default
creds, _ = default()
gc = gspread.authorize(creds)
worksheet = gc.open('Cardetails').sheet1
# get_all_values gives a list of rows.
rows = worksheet.get_all_values()
# Convert to a DataFrame and render.
import pandas as pd
df = pd.DataFrame.from_records(rows)

```

### 0.3 Data Inspection and cleaning

```

[ ]: # setting first row as headers
df.columns = df.iloc[0]
df = df.iloc[1:]
df.head()

```

```

[ ]: 0          name  year  selling_price  km_driven  fuel  \
1    Maruti Swift Dzire VDI  2014      450000    145500  Diesel
2  Skoda Rapid 1.5 TDI Ambition  2014      370000    120000  Diesel
3    Honda City 2017-2020 EXi  2006      158000    140000  Petrol

```

4	Hyundai i20 Sportz Diesel	2010	225000	127000	Diesel
5	Maruti Swift VXi BSIII	2007	130000	120000	Petrol

0	seller_type	transmission	owner	mileage	engine	max_power	\
1	Individual	Manual	First Owner	23.4 kmpl	1248 CC	74 bhp	
2	Individual	Manual	Second Owner	21.14 kmpl	1498 CC	103.52 bhp	
3	Individual	Manual	Third Owner	17.7 kmpl	1497 CC	78 bhp	
4	Individual	Manual	First Owner	23.0 kmpl	1396 CC	90 bhp	
5	Individual	Manual	First Owner	16.1 kmpl	1298 CC	88.2 bhp	

0	torque	seats
1	190Nm@ 2000rpm	5
2	250Nm@ 1500-2500rpm	5
3	12.7@ 2,700(kgm@ rpm)	5
4	22.4 kgm at 1750-2750rpm	5
5	11.5@ 4,500(kgm@ rpm)	5

```
[ ]: # size of the data
df.shape
```

```
[ ]: (8128, 13)
```

```
[ ]: # Checking for duplicates
df[df.duplicated(keep=False)]

# dropping duplicates
df.drop_duplicates(inplace=True)
```

```
[ ]: df.head()
```

[ ]: 0	name	year	selling_price	km_driven	fuel	\
1	Maruti Swift Dzire VDI	2014	450000	145500	Diesel	
2	Skoda Rapid 1.5 TDI Ambition	2014	370000	120000	Diesel	
3	Honda City 2017-2020 EXi	2006	158000	140000	Petrol	
4	Hyundai i20 Sportz Diesel	2010	225000	127000	Diesel	
5	Maruti Swift VXi BSIII	2007	130000	120000	Petrol	

0	seller_type	transmission	owner	mileage	engine	max_power	\
1	Individual	Manual	First Owner	23.4 kmpl	1248 CC	74 bhp	
2	Individual	Manual	Second Owner	21.14 kmpl	1498 CC	103.52 bhp	
3	Individual	Manual	Third Owner	17.7 kmpl	1497 CC	78 bhp	
4	Individual	Manual	First Owner	23.0 kmpl	1396 CC	90 bhp	
5	Individual	Manual	First Owner	16.1 kmpl	1298 CC	88.2 bhp	

0	torque	seats
1	190Nm@ 2000rpm	5
2	250Nm@ 1500-2500rpm	5

```

3      12.7@ 2,700(kgm@ rpm)      5
4  22.4 kgm at 1750-2750rpm      5
5      11.5@ 4,500(kgm@ rpm)      5

```

```

[ ]: # function for cleaning data
def get_clean_car_detail(car_detail):
    car_detail = car_detail.split(" ")[0]
    return car_detail

```

```

[ ]: df.name = df.name.apply(get_clean_car_detail)
df.mileage = df.mileage.apply(get_clean_car_detail)
df.engine = df.engine.apply(get_clean_car_detail)
df.max_power = df.max_power.apply(get_clean_car_detail)

```

```

[ ]: import re

def extract_torque_value(torque):
    if pd.isnull(torque):
        return None
    # Search for the first float or integer in the string
    match = re.search(r'\d+\.\d+|\d+', str(torque))
    if match:
        return float(match.group())
    return None

```

```

[ ]: df.torque = df.torque.apply(extract_torque_value)

```

```

[ ]: # setting year to type int
df.year = df.year.astype(int)

df['car_age'] = 2025 - df['year']

```

```

[ ]: # checking for missing values
df.isnull().sum()

# replacing missing values with nan
df.replace('', np.nan, inplace=True)

```

```

[ ]: df.head(2)

```

```

[ ]: 0   name  year  selling_price  km_driven   fuel  seller_type  transmission  \
1  Maruti  2014         450000   145500  Diesel  Individual        Manual
2   Skoda  2014         370000   120000  Diesel  Individual        Manual

0      owner  mileage  engine  max_power  torque  seats  car_age
1  First Owner    23.4   1248         74   190.0     5      11
2 Second Owner    21.14  1498    103.52   250.0     5      11

```

```
[ ]: df[df['max_power'].str.contains('bhp', na = False)][ 'max_power']

# Replacing bhp with nothing and replace with nan
df['max_power'] = df['max_power'].str.replace('bhp', '')
df['max_power'] = df.max_power.replace('', np.nan)
```

## 0.4 Statistical Analysis

### 0.4.1 Figure: Correlation of Numerical Variables with Target Variable (selling\_price)

```
[ ]: nums_float = ['mileage', 'engine', 'max_power', 'torque', 'selling_price',
    ↳ 'km_driven', 'seats', 'car_age']
nums_int = ['year']
df[nums_float] = df[nums_float].astype(float)
df[nums_int] = df[nums_int].astype(int)
```

```
[ ]: nums = nums_float + nums_int
len(nums)
```

```
[ ]: 9
```


```
[ ]: # Subplot
correlation_matrix = df[nums].corr()
# Focus on correlation of each feature with 'selling_price'
price_corr = correlation_matrix['selling_price']

fig, ax = plt.subplots(1, 2, figsize=(15, 6))
# Visualize the correlation with a bar plot
plt.figure(figsize=(8, 4))
price_corr.drop('selling_price').sort_values(ascending=False).plot(kind='bar',
    ↳ color='teal', ax=ax[0])
ax[0].set_title("Correlation of Features with Selling Price")
ax[0].set_ylabel("Correlation Coefficient")
ax[0].set_xlabel("Features")

# Heat Map
df[nums].corr()

# Visualizing correlation between numeric features
plt.figure(figsize=(10,4))
sns.heatmap(df[nums].corr(), annot=True, cmap = 'coolwarm', ax=ax[1])
ax[1].set_title("Correlation between Car numeric features")

plt.tight_layout()
plt.show()
```



<Figure size 800x400 with 0 Axes>

<Figure size 1000x400 with 0 Axes>

### Interpretation of the above graphs

- engine (0.69): Cars with larger engines tend to have significantly higher selling prices.
- max\_power (0.69): Cars with higher maximum power are strongly associated with higher selling prices.
- torque (0.62): Higher torque values also correlate well with higher selling prices.
- year (0.43): Newer cars tend to have higher selling prices.
- car\_age (-0.43): As cars get older, their selling price tends to decrease. This is the inverse of the correlation with year, which makes sense.
- mileage (-0.11): Higher mileage shows a weak negative correlation with selling price, suggesting that cars with more miles tend to be priced slightly lower.
- km\_driven (-0.17): Similar to mileage, it shows a weak negative correlation.
- seats (0.16): The number of seats has a weak positive correlation with selling price.

### Strong Positive Correlations Among Features (Potential Multicollinearity)

- engine, max\_power, and torque show very strong positive correlations with each other (around 0.74 - 0.8). This suggests that these features are highly related. Either engine or max\_power will be used for the model training, as their correlation with selling price is the same (0.69) and slightly greater than that of torque (0.62). Torque might be considered as an alternative if needed.
- car\_age and year have a perfect negative correlation (-1), which is expected since one is directly derived from the other. Only year will be kept for the model training to avoid

redundancy.

**In summary**, the selling price of a car is most strongly correlated with the **engine** (or **max\_power**), manufacturing **year**, and **km\_driven** (negatively). The number of **seats** will also be considered in the model despite its weak positive correlation.

The highly correlated features (**engine**, **max\_power**, **torque** and **car\_age**, **year**) will be handled by selecting one representative from each group (**engine** is preferred).

Both **mileage** and **km\_driven** measure car usage by total distance traveled. **Mileage** is in miles, while **km\_driven** is in kilometers. **km\_driven** shows a slightly stronger negative correlation with selling price (-0.17 vs. -0.11), meaning higher usage leads to a larger price drop when using **km\_driven**, assuming consistent units.

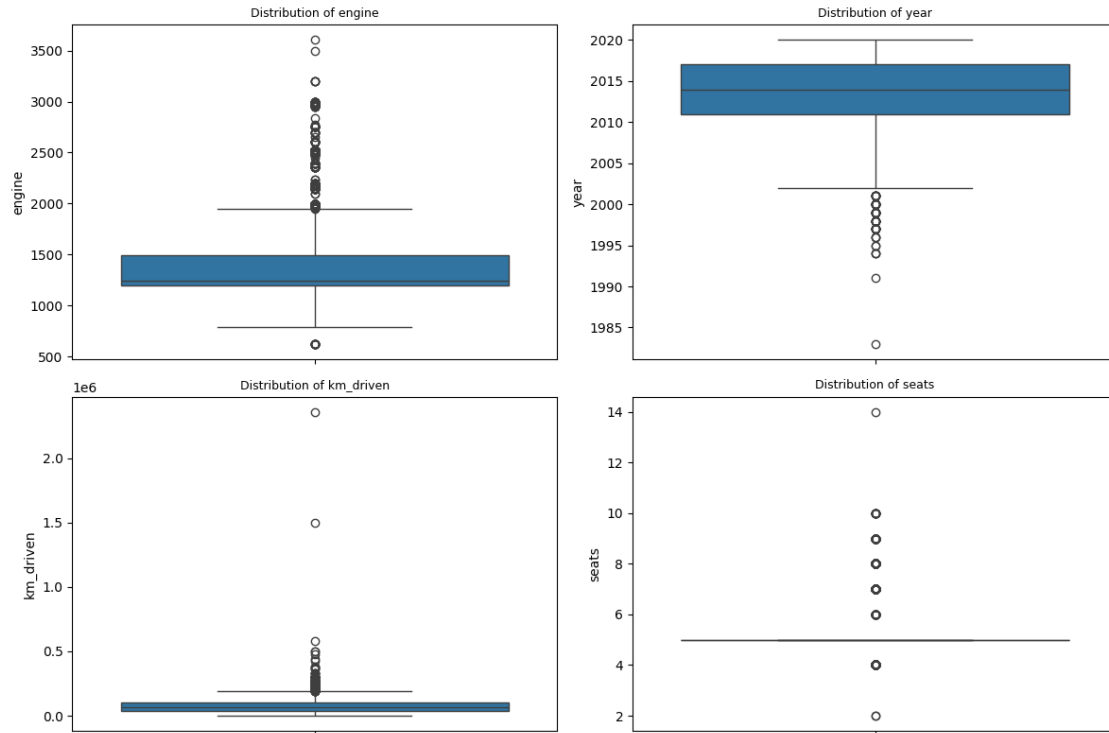
### Final Selection

- The model will include the following numerical features: **engine**, **year**, **km\_driven** and **seats**.

```
[ ]: num_features = ['engine', 'year', 'km_driven', 'seats']

# Box Plots
plt.figure(figsize=(12, 8))
for i, feature in enumerate(num_features, 1):
    plt.subplot(2, 2, i)
    sns.boxplot(y=df[feature])
    plt.title(f'Distribution of {feature}', fontsize=9)
    plt.xlabel('', fontsize=9)
plt.tight_layout()
plt.show()
```





```
[ ]: df.km_driven.max()
```

```
[ ]: 2360457.0
```

```
[ ]: # box plot showing distribution of selling price
plt.figure(figsize=(12, 5))
sns.boxplot(y=df['selling_price'])
plt.title('Distribution of Selling Price', fontsize=9)
plt.xlabel('', fontsize=9)
plt.show()
```



## Distribution of Selling Price

### Key Observations:

- **Median Inclination:** The median line within the box is noticeably closer to the lower quartile, confirming that the central tendency of the data leans towards lower selling prices.
- **Long Right Tail:** A significant number of outliers are present on the higher end of the price range, forming a long tail extending to the right. These represent a smaller proportion of cars with considerably higher selling prices.

### Impact of Distribution on Predictions:

The right-skewed distribution of the selling price is likely to influence the model's predictions in the following ways:

1. **Better Performance on Lower Prices:** Due to the higher frequency of lower-priced cars in the training data, the model is expected to perform more accurately when predicting selling prices within this dominant range.
2. **Challenges with Higher Prices:** The relatively fewer examples of high-priced cars in the training data may lead to less accurate and more variable predictions for these vehicles.
3. **Potential for Underprediction:** Models trained on right-skewed data can sometimes exhibit a tendency to underpredict values in the longer, higher-value tail.

Understanding this distribution is crucial for interpreting the model's performance metrics and the "Predicted vs Actual Prices" plot later.

```
[ ]: df.describe()
```

```
[ ]: 0          year  selling_price  km_driven  mileage  engine \
count  6926.000000  6.926000e+03  6.926000e+03  6718.00000  6718.000000
mean    2013.420300  5.172707e+05  7.399568e+04  19.46531  1430.891337
std       4.078286  5.197670e+05  5.835810e+04  4.04915  493.493277
min     1983.000000  2.999900e+04  1.000000e+00  0.00000  624.000000
```

25%	2011.000000	2.500000e+05	4.000000e+04	16.80000	1197.000000
50%	2014.000000	4.000000e+05	7.000000e+04	19.44000	1248.000000
75%	2017.000000	6.335000e+05	1.000000e+05	22.50000	1498.000000
max	2020.000000	1.000000e+07	2.360457e+06	42.00000	3604.000000


	max_power	torque	seats	car_age
count	6720.000000	6717.000000	6718.000000	6926.000000
mean	87.726919	160.854853	5.434653	11.579700
std	31.771619	91.630280	0.984230	4.078286
min	0.000000	4.800000	2.000000	5.000000
25%	67.100000	96.000000	5.000000	8.000000
50%	81.830000	146.000000	5.000000	11.000000
75%	100.000000	200.000000	5.000000	14.000000
max	400.000000	789.000000	14.000000	42.000000

```
[ ]: cat = ['name', 'fuel', 'seller_type', 'transmission', 'owner']
      target = ['selling_price']
      cat_target = cat + target
      df[cat_target].head()
```

```
[ ]: 0    name    fuel seller_type transmission    owner    selling_price
     1  Maruti  Diesel  Individual    Manual    First Owner    450000.0
     2  Skoda  Diesel  Individual    Manual    Second Owner    370000.0
     3  Honda  Petrol  Individual    Manual    Third Owner    158000.0
     4  Hyundai Diesel  Individual    Manual    First Owner    225000.0
     5  Maruti  Petrol  Individual    Manual    First Owner    130000.0
```

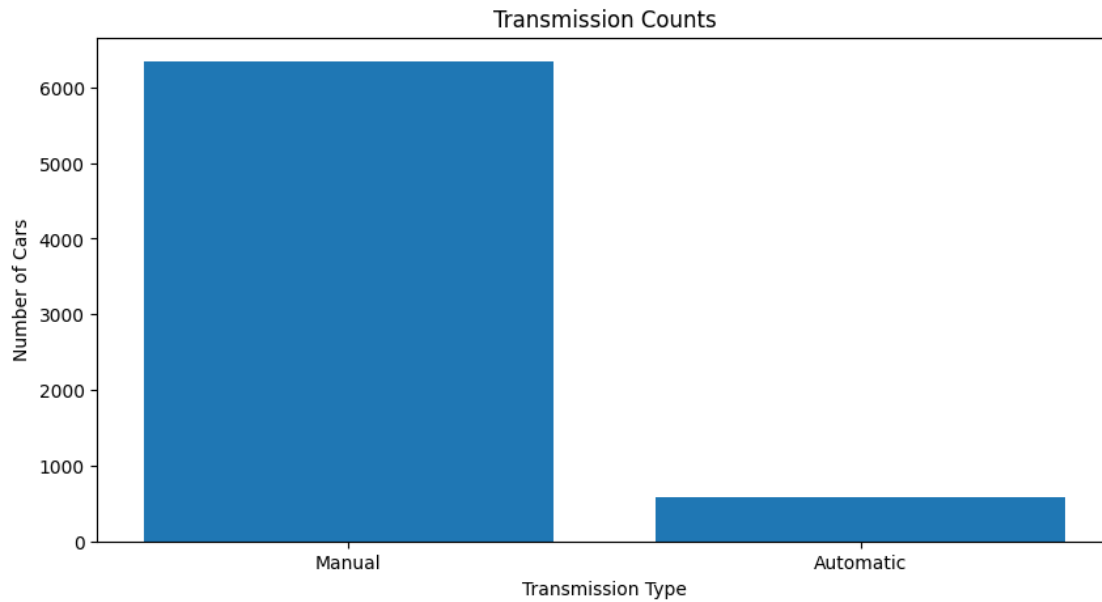
```
[ ]: # counts for Owner
      owner_counts = df['owner'].value_counts()

      # Plot
      plt.figure(figsize=(10,5))
      plt.bar(owner_counts.index, owner_counts.values)
      plt.title('Owner Counts')
      plt.xlabel('Owner Type')
      plt.ylabel('Number of Cars')
      plt.show()
```



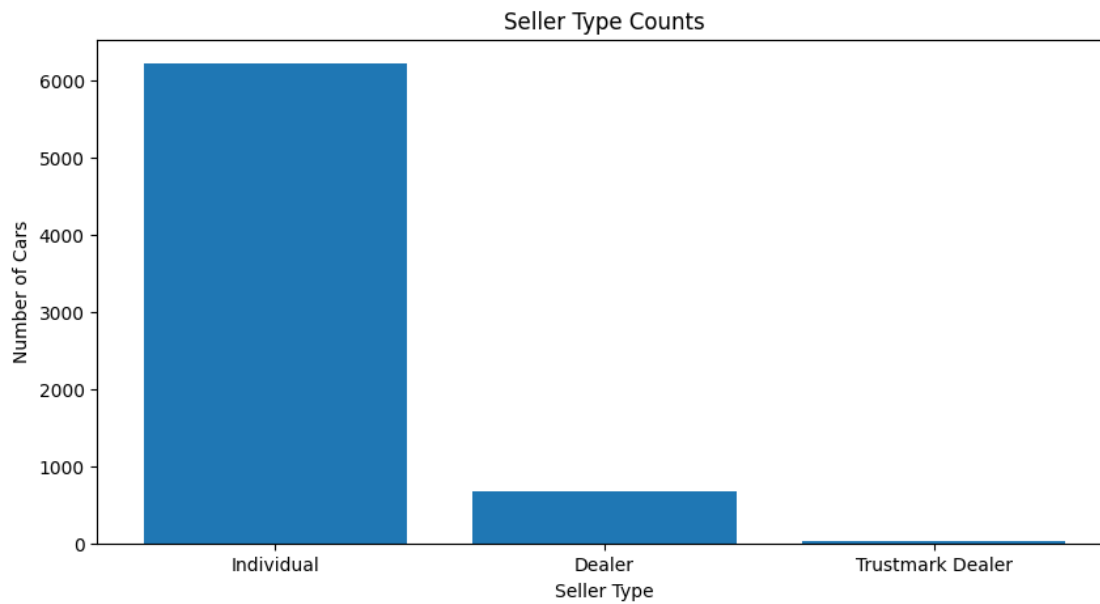
```
[ ]: # Transmission counts
      transmission_counts = df['transmission'].value_counts()

      # Plot
      plt.figure(figsize=(10,5))
      plt.bar(transmission_counts.index, transmission_counts.values)
      plt.title('Transmission Counts')
      plt.xlabel('Transmission Type')
      plt.ylabel('Number of Cars')
      plt.show()
```



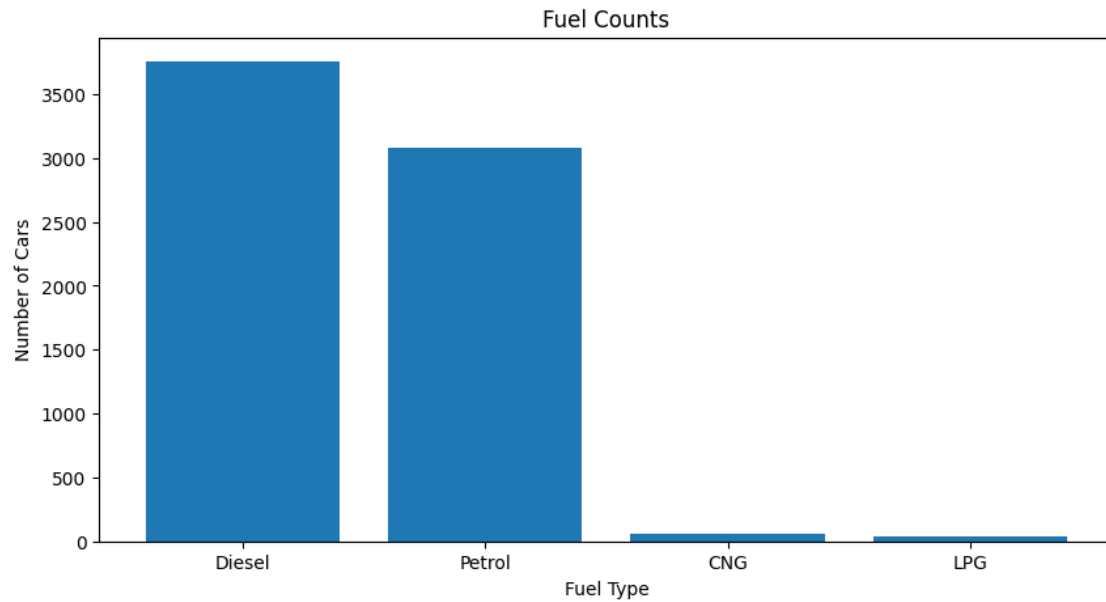
```
[ ]: # Counts (seller_type)
seller_counts = df['seller_type'].value_counts()

# Plot
plt.figure(figsize=(10,5))
plt.bar(seller_counts.index, seller_counts.values)
plt.title('Seller Type Counts')
plt.xlabel('Seller Type')
plt.ylabel('Number of Cars')
plt.show()
```



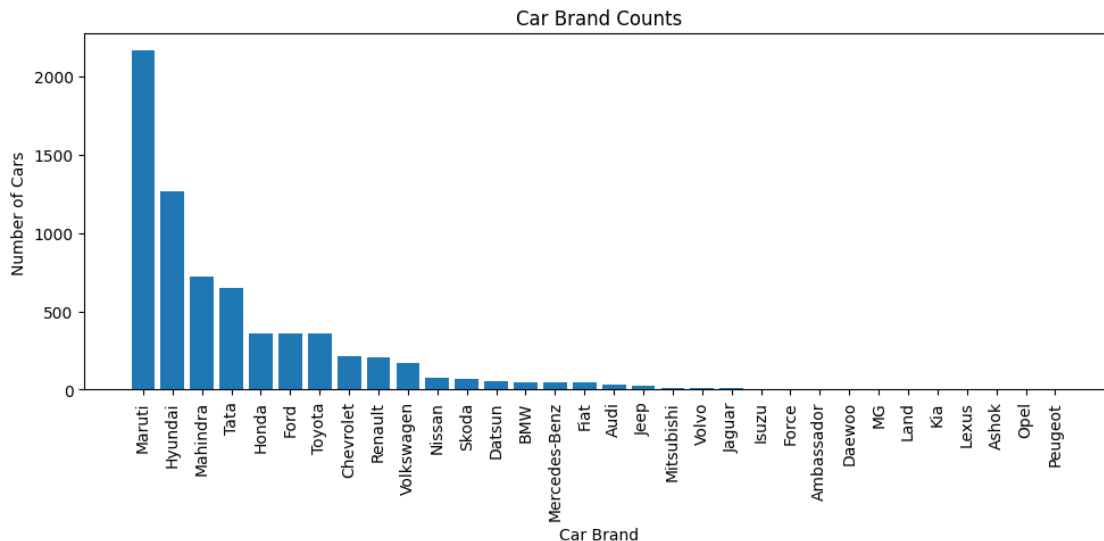
```
[ ]: # Counts for Fuel
fuel_counts = df['fuel'].value_counts()

# Plot
plt.figure(figsize=(10,5))
plt.bar(fuel_counts.index, fuel_counts.values)
plt.title('Fuel Counts')
plt.xlabel('Fuel Type')
plt.ylabel('Number of Cars')
plt.show()
```



```
[ ]: # Get counts of each car brand
name_counts = df['name'].value_counts()

# Plot
plt.figure(figsize=(10,5))
plt.bar(name_counts.index, name_counts.values)
plt.xticks(rotation=90) # Rotate x-axis labels for readability
plt.title('Car Brand Counts')
plt.xlabel('Car Brand')
plt.ylabel('Number of Cars')
plt.tight_layout()
plt.show()
```



```
[ ]: anova_results = []
      for col in cat:
          # performing welch anova
          formula = f'selling_price ~ C({col})'
          model = ols(formula, data=df).fit()
          anova_table = sm.stats.anova_lm(model, typ=2, robust='hc3')
          anova_table = anova_table[anova_table.index != 'Residual']
          anova_results.append(anova_table)

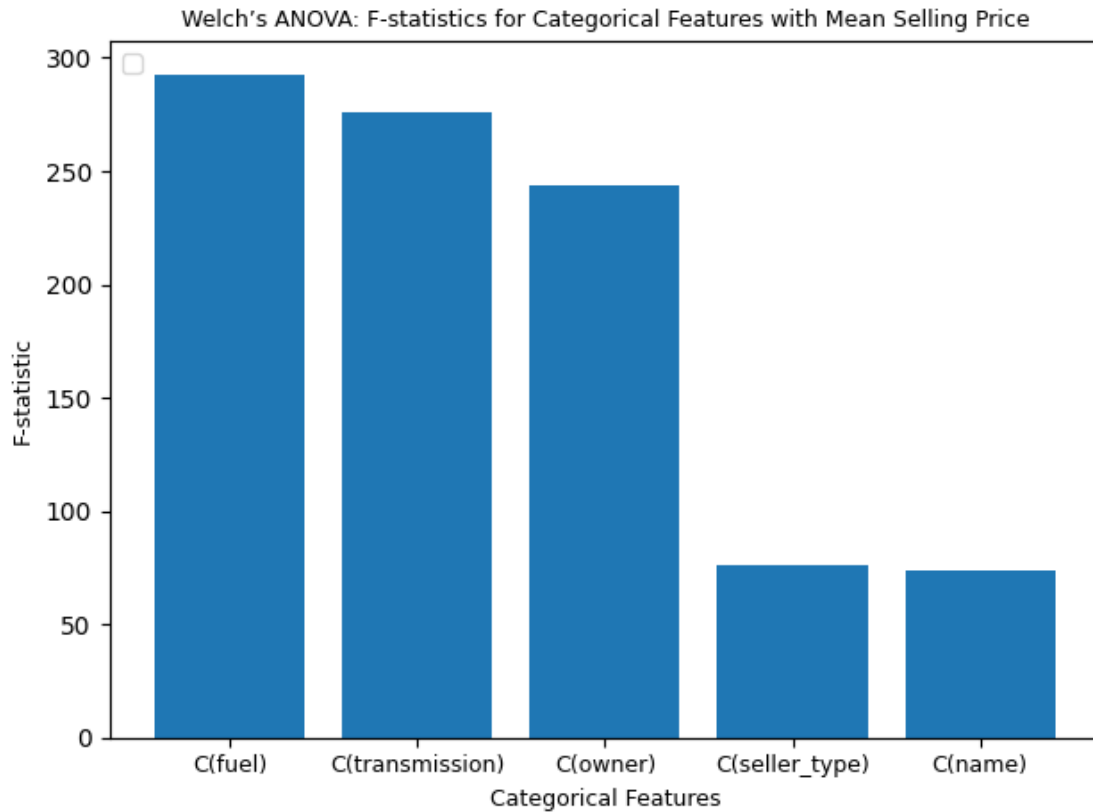
      anova_results_df = pd.concat(anova_results, axis=0)
      anova_results_df = anova_results_df.sort_values(by='F', ascending=False)
      anova_results_df
```

```
[ ]:
      sum_sq    df      F      PR(>F)
C(fuel)      2.214430e+14  3.0  292.567986  8.080008e-179
C(transmission)  5.860624e+13  1.0  275.983306  8.374010e-61
C(owner)      2.361284e+14  4.0  243.779556  3.437094e-196
C(seller_type)  3.844585e+13  2.0   76.369348  1.562395e-33
C(name)      3.289485e+14 31.0   73.405500  0.000000e+00
```

```
[ ]: plt.bar(anova_results_df.index, anova_results_df['F'])
      plt.xlabel('Categorical Features', fontsize = 9)
      plt.ylabel('F-statistic', fontsize = 9)
      plt.title('Welch's ANOVA: F-statistics for Categorical Features with Mean_
        ↳ Selling Price', fontsize = 9)
      plt.xticks(fontsize = 9)
      plt.legend(loc='upper left')
      plt.tight_layout()
```



```
plt.show()
```



### Understanding Welch's ANOVA F-statistic

Welch's ANOVA is a statistical test used to compare the means of two or more groups when the population variances are not assumed to be equal (unlike the standard ANOVA). Here, the F-statistic indicates the ratio of variance between group means to average variance within group means.

A higher F-statistic generally suggests a greater difference in the mean selling price across the different categories in that group. In other words, the feature is likely a better predictor of selling price.

1. **C(fuel)**: This feature has the highest F-statistic. This strongly suggests that the type of fuel a car uses has a substantial impact on its selling price. There are likely statistically significant differences in the average selling prices of cars with different fuel types.
2. **C(transmission)**: The transmission type also shows a relatively high F-statistic, although lower than fuel. This indicates that the type of transmission (e.g., manual, automatic) is also a significant factor influencing the selling price. We can infer that cars with different transmission types tend to have different average selling prices.
3. **C(owner)**: The F-statistic for the number of previous owners is moderate. This suggests that the number of owners has a noticeable, but less pronounced than fuel or transmission,

effect on the selling price. Cars with fewer previous owners likely have different average selling prices compared to those with more owners.

4. **C(seller\_type)**: While seller\_type shows statistical significance with a notable F-statistic, it has a lower impact on selling price compared to features like fuel type, transmission, and number of previous owners. This variable reflects the context of the transaction rather than characteristics of the car. Since the model is intended for use by car sellers to estimate prices, including seller\_type could introduce data leakage or bias. Therefore, despite statistical relevance, seller\_type is intentionally excluded to ensure fairness, robustness, and cleaner interpretation of results.
5. **C(name)**: The F-statistic for the car's name (make and model) is also relatively low. The F-statistic here might be lower because this is a high-cardinality categorical feature (many unique car names) and is highly imbalanced as seen from the graph above (Car brands count). These characteristics dilute the overall variance captured by ANOVA, even if individual brands have meaningful differences in selling price. The make and model of a car often encapsulate key pricing determinants such as brand reputation, performance, and market segment. Therefore, despite statistical challenges, the name feature holds substantial predictive value and is to be included in the modeling phase with appropriate encoding techniques.

## Final Selection

- The model will include the following categorical features: **fuel**, **transmission**, **owner** and **name**.

## 0.5 Model Training

```
[ ]: df_model = df.drop(columns = ['seller_type', 'mileage', 'max_power', 'torque', 'car_age'])
```

```
[ ]: df_model.head()
```

```
[ ]: 0    name  year  selling_price  km_driven  fuel transmission  owner \
1  Maruti  2014    450000.0    145500.0  Diesel      Manual  First Owner
2  Skoda   2014    370000.0    120000.0  Diesel      Manual  Second Owner
3  Honda   2006    158000.0    140000.0  Petrol      Manual  Third Owner
4  Hyundai 2010    225000.0    127000.0  Diesel      Manual  First Owner
5  Maruti  2007    130000.0    120000.0  Petrol      Manual  First Owner

0  engine  seats
1  1248.0    5.0
2  1498.0    5.0
3  1497.0    5.0
4  1396.0    5.0
5  1298.0    5.0
```

```
[ ]: # defining features and the target
x = df_model.drop(columns = ['selling_price'])
y = df_model['selling_price']
```

```
[ ]: # selecting the numerical features in a datatype int and float
num_model_features = ['year', 'km_driven', 'engine', 'seats']
x[num_model_features].head()
```

```
[ ]: 0  year  km_driven  engine  seats
1  2014   145500.0   1248.0    5.0
2  2014   120000.0   1498.0    5.0
3  2006   140000.0   1497.0    5.0
4  2010   127000.0   1396.0    5.0
5  2007   120000.0   1298.0    5.0
```

```
[ ]: cat_model_features = ['name', 'fuel', 'transmission', 'owner']
x[cat_model_features].head()
```

```
[ ]: 0    name    fuel transmission    owner
1  Maruti  Diesel      Manual  First Owner
2   Skoda  Diesel      Manual  Second Owner
3   Honda  Petrol      Manual  Third Owner
4 Hyundai  Diesel      Manual  First Owner
5  Maruti  Petrol      Manual  First Owner
```

```
[ ]: # Features pipeline
num_pipeline = Pipeline([
    ('imputer', SimpleImputer(strategy='mean')),
    ('scaler', StandardScaler())
])

cat_pipeline = Pipeline([
    ('onehot', OneHotEncoder(handle_unknown='ignore'))
])
```

```
[ ]: # Column transformer
preprocessor = ColumnTransformer([
    ('num', num_pipeline, num_model_features),
    ('cat', cat_pipeline, cat_model_features)
])

# display the preprocessor
preprocessor
```

```
[ ]: ColumnTransformer(transformers=[('num',
                                     Pipeline(steps=[('imputer', SimpleImputer()),
                                                         ('scaler', StandardScaler())]),
                                     ['year', 'km_driven', 'engine', 'seats']),
                                    ('cat',
                                     Pipeline(steps=[('onehot',
                                                         OneHotEncoder(handle_unknown='ignore'))]),
```

```
[ 'name', 'fuel', 'transmission', 'owner' ] ] ] )
```

```
[ ]: # splitting the data
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2,
random_state=42)
```

```
[ ]: x_test.head()
```

```
[ ]: 0      name  year  km_driven  fuel transmission      owner  engine \
8078  Toyota  2009   250000.0  Diesel      Manual  First Owner  2494.0
4095  Mahindra 2015    39000.0  Diesel      Manual  First Owner  2179.0
6494   Ford  2018    90000.0  Diesel      Manual  First Owner  1498.0
4340  Maruti  2008   100000.0  Petrol      Manual  Third Owner  1061.0
2536   Honda  2018    40000.0  Diesel      Manual  First Owner  1498.0

0      seats
8078    7.0
4095    7.0
6494    5.0
4340    5.0
2536    5.0
```

### 0.5.1 Random Forest Model

```
[ ]: # Rf pipeline
rf_pipeline = Pipeline([
    ('preprocessor', preprocessor),
    ('rf', RandomForestRegressor(n_estimators=100, random_state=42))
])

# displaying the reg pipeline
rf_pipeline
```

```
[ ]: Pipeline(steps=[('preprocessor',
                      ColumnTransformer(transformers=[('num',
                                                         Pipeline(steps=[('imputer',
                                                                              SimpleImputer()),
                                                                              ('scaler',
                                                                              StandardScaler())])),
                                                         ('cat',
                                                         Pipeline(steps=[('onehot',
                                                                              OneHotEncoder(handle_unknown='ignore'))])),
                                                         ['year', 'km_driven',
                                                                              'engine', 'seats']),
                                                         ('name', 'fuel',
                                                                              'transmission',
                                                                              'owner'])])),
```

```
(('rf', RandomForestRegressor(random_state=42)))]])
```

```
[ ]: # Training the data using Random forest Regressor
rf_model = rf_pipeline.fit(x_train, y_train)
rf_model
```

```
[ ]: Pipeline(steps=[('preprocessor',
                      ColumnTransformer(transformers=[('num',
                                                         Pipeline(steps=[('imputer',
                                                                              SimpleImputer()),
                                                                              ('scaler',
                                                                              StandardScaler()))],
                                                         ['year', 'km_driven',
                                                         'engine', 'seats']),
                                                         ('cat',
                                                         Pipeline(steps=[('onehot',
                                                                              OneHotEncoder(handle_unknown='ignore'))]),
                                                         ['name', 'fuel',
                                                         'transmission',
                                                         'owner'])])),
                      ('rf', RandomForestRegressor(random_state=42)))]])
```

```
[ ]: # Predict
y_pred_rf = rf_model.predict(x_test)
```

```
[ ]: print("MAE:", round(mean_absolute_error(y_test, y_pred_rf),2))
print("MSE:", round(mean_squared_error(y_test, y_pred_rf),2))
print("RMSE:", round(np.sqrt(mean_squared_error(y_test, y_pred_rf)),2))
print("R2 Score:", r2_score(y_test, y_pred_rf))
```

MAE: 91316.73

MSE: 39347158917.14

RMSE: 198361.18

R<sup>2</sup> Score: 0.8205960840365168

## 0.5.2 Gradient Boosting and GridSearchCV

```
[ ]: # Hyperparameters
param_grid = {
    'gsb__n_estimators': [100, 200],
    'gsb__learning_rate': [0.01, 0.1, 0.2],
    'gsb__max_depth': [3, 4, 5]
}
```

```
[ ]: # gsb pipeline
gb_pipeline = Pipeline([
    ('preprocessor', preprocessor),
```

```

        ('gsb', GradientBoostingRegressor(random_state=42))
    ])

    # gsb parameters
    grid_search = GridSearchCV(gb_pipeline, param_grid, cv=5,
                               scoring='neg_mean_squared_error', n_jobs=-1)
    grid_search

```

```

[ ]: GridSearchCV(cv=5,
                  estimator=Pipeline(steps=[('preprocessor',
                                             ColumnTransformer(transformers=[('num',
Pipeline(steps=[('imputer',
                  SimpleImputer()),
                  ('scaler',
                    StandardScaler()))]),
                ['year',
                 'km_driven',
                 'engine',
                 'seats']),
                ('cat',
Pipeline(steps=[('onehot',
                  OneHotEncoder(handle_unknown='ignore'))]),
                ['name',
                 'fuel',
                 'transmission',
                 'owner'])])),
                  ('gsb',
GradientBoostingRegressor(random_state=42))),
        n_jobs=-1,
        param_grid={'gsb__learning_rate': [0.01, 0.1, 0.2],
                     'gsb__max_depth': [3, 4, 5],
                     'gsb__n_estimators': [100, 200]},
        scoring='neg_mean_squared_error')

```

```

[ ]: # Training
gs_model = grid_search.fit(x_train, y_train)

```

```

[ ]: # Best params
best_params = gs_model.best_params_
best_params

```

```

[ ]: {'gsb__learning_rate': 0.2, 'gsb__max_depth': 5, 'gsb__n_estimators': 100}

```

```

[ ]: # best model
best_model = gs_model.best_estimator_

```

```
[ ]: # predicting
y_pred_gb = best_model.predict(x_test)
```

```
[ ]: print("MAE:", round(mean_absolute_error(y_test, y_pred_gb),2))
print("MSE:", round(mean_squared_error(y_test, y_pred_gb),2))
print("RMSE:", round(np.sqrt(mean_squared_error(y_test, y_pred_gb)),2))
print("R2 Score:", r2_score(y_test, y_pred_gb))
```

MAE: 88424.5  
MSE: 36732389083.97  
RMSE: 191656.96  
R<sup>2</sup> Score: 0.8325181633002938

### 0.5.3 Xgboost

```
[ ]: xgb_pipeline = Pipeline([
    ('preprocessor', preprocessor),
    ('xgb', XGBRegressor(
        objective='reg:squarederror',
        eval_metric='rmse',
        random_state=42
    ))
])

# display the pipeline
xgb_pipeline
```

```
[ ]: Pipeline(steps=[('preprocessor',
    ColumnTransformer(transformers=[('num',
        Pipeline(steps=[('imputer',
            SimpleImputer()),
            ('scaler',
                StandardScaler())])),
        [ 'year', 'km_driven',
          'engine', 'seats']),
        ('cat',
            Pipeline(steps=[('onehot',
                OneHotEncoder(handle_unknown='ignore'))]),
            [ 'name', 'fuel',
              'transmission',
              'owner'])])),
    ('xgb',
        XGBRegressor(base_score=None, booster=None,
            feature_types=None, gamma=None, grow_policy=None,
            importance_type=None,
            interaction_constraints=None, learning_rate=None,
            max_bin=None, max_cat_threshold=None,
```

```

max_cat_to_onehot=None, max_delta_step=None,
max_depth=None, max_leaves=None,
min_child_weight=None, missing=nan,
monotone_constraints=None, multi_strategy=None,
n_estimators=None, n_jobs=None,
num_parallel_tree=None, random_state=42, ...)))

```

```

[ ]: # Train the model
xgb_model = xgb_pipeline.fit(x_train, y_train)

```

```

[ ]: # Predicting
y_pred_xgb = xgb_model.predict(x_test)

```

```

[ ]: print("MAE:", round(mean_absolute_error(y_test, y_pred_xgb),2))
print("MSE:", round(mean_squared_error(y_test, y_pred_xgb),2))
print("RMSE:", round(np.sqrt(mean_squared_error(y_test, y_pred_xgb)),2))
print("R2 Score:", r2_score(y_test, y_pred_xgb))

```

MAE: 87513.79

MSE: 32608928997.74

RMSE: 180579.43

R<sup>2</sup> Score: 0.8513191366652719


```

[ ]: # Model Performance Visualization

plt.figure(figsize=(8, 6))
sns.scatterplot(x=y_test, y=y_pred_xgb, alpha=0.7)
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], color='red', linestyle='--')
plt.xlabel('Actual Price')
plt.ylabel('Predicted Price')
plt.title('Predicted vs Actual Prices')
plt.grid(True)
plt.show()

```





### Interpretation

- **X-axis (Actual Price):** This represents the true selling prices of the cars.
- **Y-axis (Predicted Price):** This represents the selling prices predicted by `xgb_model`.
- **Red Dashed Line (Ideal Prediction):** This diagonal line represents a scenario where the predicted price is exactly equal to the actual price. If all the points fell perfectly on this line, the model would be making perfect predictions.
- **Blue Scatter Points:** Each blue dot represents a single car. Its position on the graph shows the actual selling price (x-coordinate) and the price predicted (y-coordinate).
- **Clustering Around the Red Line:** Notice that a significant number of the blue points are clustered relatively close to the red dashed line, especially for cars with lower actual selling prices (below roughly 2 million). This indicates that the model is generally performing well for cars in this price range.
- **Scatter at Higher Prices:** As the actual selling price increases (moving to the right on the x-axis), the scatter of the blue points tends to become wider around the red line. This suggests that the model's predictions become less precise for higher-priced cars. The errors (the vertical distance between a point and the red line) are larger for these more expensive vehicles.

- Points above the red line represent cases where your model overpredicted the selling price (the predicted price is higher than the actual price).
- Points below the red line represent cases where your model underpredicted the selling price (the predicted price is lower than the actual price).
- In conclusion, the model is reasonably good at predicting the selling prices of cars, particularly those in the lower to mid-price range. However, it exhibits a higher degree of error and less precision when predicting the prices of more expensive vehicles.

#### 0.5.4 Saving the xgb model

```
[ ]: # displaying
xgb_model
```

```
[ ]: Pipeline(steps=[('preprocessor',
                      ColumnTransformer(transformers=[('num',
                                                         Pipeline(steps=[('imputer',
                                                                              SimpleImputer()),
                                                                              ('scaler',
                                                                              StandardScaler()))]),
                                                         ['year', 'km_driven',
                                                                              'engine', 'seats']),
                                                         ('cat',
                                                         OneHotEncoder(handle_unknown='ignore'))])),
                ('xgb',
                 XGBRegressor(base_score=None, booster=None,
                               feature_types=None, gamma=None, grow_policy=None,
                               importance_type=None,
                               interaction_constraints=None, learning_rate=None,
                               max_bin=None, max_cat_threshold=None,
                               max_cat_to_onehot=None, max_delta_step=None,
                               max_depth=None, max_leaves=None,
                               min_child_weight=None, missing=nan,
                               monotone_constraints=None, multi_strategy=None,
                               n_estimators=None, n_jobs=None,
                               num_parallel_tree=None, random_state=42, ...)))
```

```
[ ]: # saving the xgb model
joblib.dump(xgb_model, 'xgb_model.pkl')
```

```
[ ]: ['xgb_model.pkl']
```


```
[ ]: x_train.head()
```

```
[ ]: 0      name  year  km_driven  fuel  transmission  owner  engine  \
5587   Ford   2018   25000.0  Diesel      Manual   First Owner  1498.0
3469   Maruti 2010   170000.0 Petrol      Manual   Second Owner   NaN
4936   Hyundai 2011    75500.0 Diesel      Manual   Second Owner  1396.0
2486   Maruti 2009   138000.0 Diesel      Manual   Second Owner  1248.0
4492   Hyundai 2003   200000.0 Diesel      Manual   Second Owner  1493.0

0      seats
5587    5.0
3469    NaN
4936    5.0
2486    5.0
4492    5.0
```

## 1 Model 2 - Removed extreme outliers

```
[ ]: # box plot showing distribution of selling price
plt.figure(figsize=(12, 5))
sns.boxplot(y=df['selling_price'])
plt.title('Distribution of Selling Price', fontsize=9)
plt.xlabel('', fontsize=9)
plt.show()
```



Removing data points beyond the upper whisker (i.e., high outliers).

```
[ ]: # q1 and q3
q1 = df["selling_price"].quantile(0.25)
q3 = df['selling_price'].quantile(0.75)
```

```
# inter-quantile range
iqr = q3 - q1

# upper whisker
upper_whisker = q3 + 1.5 * iqr

# removing the outliers
df_no_outliers = df[df['selling_price'] < upper_whisker]
```

```
[ ]: df.shape
```

```
[ ]: (6926, 14)
```

```
[ ]: print(df.shape)
print(df_no_outliers.shape)

df_rows = df.shape[0]
df_no_outliers_srows = df_no_outliers.shape[0]


lost_data = df_rows - df_no_outliers_srows
lost_data_percent = round((lost_data / df_rows * 100), 2)
print(f'Data lost is {lost_data_percent}%')
```

```
(6926, 14)
```

```
(6598, 14)
```

```
Data lost is 4.74%
```

```
[ ]: # distribution of selling price without outliers
plt.figure(figsize=(12, 5))
sns.boxplot(y=df_no_outliers['selling_price'])
plt.title('Distribution of Selling Price', fontsize=9)
plt.xlabel('', fontsize=9)
plt.show()
```



```
[ ]: df_no_outliers.drop(columns = ['seller_type', 'mileage', 'max_power', 'torque', 'car_age'], inplace=True)
df_no_outliers.head()
```

```
[ ]: 0    name  year  selling_price  km_driven  fuel transmission  owner \
1   Maruti  2014    450000.0    145500.0  Diesel      Manual  First Owner
2    Skoda  2014    370000.0    120000.0  Diesel      Manual  Second Owner
3   Honda  2006    158000.0    140000.0  Petrol      Manual  Third Owner
4  Hyundai  2010    225000.0    127000.0  Diesel      Manual  First Owner
5   Maruti  2007    130000.0    120000.0  Petrol      Manual  First Owner

0  engine  seats
1  1248.0    5.0
2  1498.0    5.0
3  1497.0    5.0
4  1396.0    5.0
5  1298.0    5.0
```

```
[ ]: # Features and target
x_no_outliers = df_no_outliers.drop(columns = ['selling_price'])
y_no_outliers = df_no_outliers['selling_price']
```

```
[ ]: # splitting the not outliers data
x_train_no_outliers, x_test_no_outliers, y_train_no_outliers, y_test_no_outliers = train_test_split(x_no_outliers, y_no_outliers, test_size = 0.2, random_state=42)
```

```
[ ]: xgb_pipeline_no_outliers = Pipeline([
    ('preprocessor', preprocessor),
    ('xgb', XGBRegressor(
        objective='reg:squarederror',
        eval_metric='rmse',
        random_state=42
    ))
])
```


```
[ ]: # training the no outliers data
xgb_model_no_outliers = xgb_pipeline_no_outliers.fit(x_train_no_outliers, y_train_no_outliers)
```

```
[ ]: # Predictions on no outliers
y_pred_xgb_no_outliers = xgb_model_no_outliers.predict(x_test_no_outliers)
```

```
[ ]: print("MAE:", round(mean_absolute_error(y_test_no_outliers,
    ↪ y_pred_xgb_no_outliers),2))
print("MSE:", round(mean_squared_error(y_test_no_outliers,
    ↪ y_pred_xgb_no_outliers),2))
print("RMSE:", round(np.sqrt(mean_squared_error(y_test_no_outliers,
    ↪ y_pred_xgb_no_outliers)),2))
print("R2 Score:", round(r2_score(y_test_no_outliers, y_pred_xgb_no_outliers),
    ↪ 2))
```

MAE: 65224.45  
MSE: 8681822530.02  
RMSE: 93176.3  
R<sup>2</sup> Score: 0.86

```
[ ]: # Graphical Representation of the model
plt.figure(figsize=(8, 6))
sns.scatterplot(x=y_test_no_outliers, y=y_pred_xgb_no_outliers, alpha=0.7)
# Perfect prediction line
plt.plot([y_test_no_outliers.min(), y_test_no_outliers.max()],
    ↪ [y_test_no_outliers.min(), y_test_no_outliers.max()], color='red',
    ↪ linestyle='--')
plt.xlabel('Actual Price')
plt.ylabel('Predicted Price')
plt.title('Predicted vs Actual Prices')
plt.grid(True)
plt.show()
```



```
[ ]: # Results for each model
results = {
    'Model': ['XGBoost', 'XGBoost (no outliers)', 'Gradient Boosting', 'Random_
↳Forest'],
    'MAE': [
        round(mean_absolute_error(y_test, y_pred_xgb), 2),
        round(mean_absolute_error(y_test_no_outliers, y_pred_xgb_no_outliers),
↳2),
        round(mean_absolute_error(y_test, y_pred_gb), 2),
        round(mean_absolute_error(y_test, y_pred_rf), 2)
    ],
    'MSE': [
        round(mean_squared_error(y_test, y_pred_xgb), 2),
        round(mean_squared_error(y_test_no_outliers, y_pred_xgb_no_outliers),
↳2),
        round(mean_squared_error(y_test, y_pred_gb), 2),
        round(mean_squared_error(y_test, y_pred_rf), 2)
    ],
    'RMSE': [
```

```

        round(np.sqrt(mean_squared_error(y_test, y_pred_xgb)), 2),
        round(np.sqrt(mean_squared_error(y_test_no_outliers,
        ↪y_pred_xgb_no_outliers)), 2),
        round(np.sqrt(mean_squared_error(y_test, y_pred_gb)), 2),
        round(np.sqrt(mean_squared_error(y_test, y_pred_rf)), 2)
    ],
    'R² Score': [
        round(r2_score(y_test, y_pred_xgb), 4),
        round(r2_score(y_test_no_outliers, y_pred_xgb_no_outliers), 4),
        round(r2_score(y_test, y_pred_gb), 4),
        round(r2_score(y_test, y_pred_rf), 4)
    ]
}
# data frame
results_df = pd.DataFrame(results)

```

```

[ ]: results_df = results_df.sort_values(by='R² Score', ascending=False)
results_df

```

```

[ ]:

```

	Model	MAE	MSE	RMSE	R² Score
1	XGBoost (no outliers)	65,224.45	8,681,822,530.02	93,176.30	0.86
0	XGBoost	87,513.79	32,608,928,997.74	180,579.43	0.85
2	Gradient Boosting	88,424.50	36,732,389,083.97	191,656.96	0.83
3	Random Forest	91,316.73	39,347,158,917.14	198,361.18	0.82

### Explanation (Model Evaluation)

- Model: The name of the machine learning model.
- MAE: Mean Absolute Error. This measures the average absolute difference between the predicted selling price and the actual selling price. Lower is better.
- MSE: Mean Squared Error. This measures the average squared difference between the predicted and actual selling prices. Lower is better.
- RMSE: Root Mean Squared Error. This is the square root of the MSE. It has the same units as the target variable (selling price), making it more interpretable than MSE. Lower is better.

---

### Performance Comparison:

- XGBoost appears to be the best-performing model based on these metrics. It has the lowest MAE, MSE, and RMSE, and the highest R² score. This indicates that, on average, its predictions are closest to the actual selling prices, with the least amount of error and the highest proportion of variance explained.
- Gradient Boosting performed second best, with MAE and RMSE values slightly higher than XGBoost but better than Random forest. Its R² score is also higher than Random forest.
- Random forest shows the weakest performance among the three, with the highest MAE, MSE, and RMSE, and the lowest R² score.



---

## XGBoost Vs XGBoost (No Outliers)

1. **XGBoost (No Outliers):** Shows Noticeable Performance Gains. Removing outliers from the training data led to a clear and measurable improvement in all evaluation metrics for the XGBoost model.
2. **Drastic Drop in RMSE Reflects Higher Accuracy:** The significant reduction in RMSE indicates that the model's predictions are now much closer to the actual selling prices. Notice that the Mean Squared Error (MSE) also dropped heavily from 32.6 billion to just 8.7 billion. Outliers previously distorted the model's learning process, but filtering them out allowed the model to better capture the general patterns in the data and make more reliable predictions.

### 1.0.1 Saving the xgb\_model\_no\_outliers

```
[ ]: # saving the model
joblib.dump(xgb_model_no_outliers, 'xgb_model_no_outliers.pkl')

[ ]: ['xgb_model_no_outliers.pkl']
```

### 1.0.2 Deployment of Streamlit App

The app was developed and tested in VSCode, then deployed using Streamlit Cloud for easy access and sharing.

```
[ ]: import pandas as pd
import streamlit as st
import xgboost as xgb
from xgboost import XGBRegressor
import joblib

st.set_page_config(layout="wide")

model = joblib.load('xgb_model.pkl')

st.header("Car Price Prediction (XGBRegressor)")
df = pd.read_csv(r"C:\Users\PC\Desktop\car_model\Cardetails - Cardetails.csv")

# function for cleaning data
def get_clean_car_detail(car_detail):
    car_detail = car_detail.split(" ")[0]
    return car_detail
df.name = df.name.apply(get_clean_car_detail)

# creating columns
col1, col2 = st.columns(2)
```

```

# Car form
with col1:
    name = st.selectbox('Select Car Brand', df['name'].unique())
    year = st.selectbox('Select Car Year', df['year'].unique())
    km_driven = st.slider("Km_driven", 0, 2500000)
    fuel = st.selectbox('Select Fuel Type', df['fuel'].unique())

with col2:
    transmission = st.selectbox('Transmission type', df['transmission'].unique())
    owner = st.selectbox('Owner', df['owner'].unique())
    engine = st.slider("Engine Size (cc)", 200, 10000)
    seats = st.slider("Number of Seats", 2, 15)

# after getting information from user
if st.button("Predict"):
    input_data_model = pd.DataFrame(
        [[name, year, km_driven, fuel, transmission,
          owner, engine, seats]],
        columns = ['name', 'year', 'km_driven', 'fuel', 'transmission',
                  'owner', 'engine', 'seats']
    )

    # st.write(input_data_model)

    # Passing values to the model to predict
    car_price = model.predict(input_data_model)

    st.markdown('Car price is {:.0f}'.format(car_price[0]))

```